

Deep Reinforcement Learning Expert Nanodegree

João Pedro Megid Carrilho

Report

Project 3: Collaboration and Competition

Introduction

In this project an Multi-Agent Actor-Critic Method of Reinforcement Learning is utilized for training a agent in a 3D simulated environment. The agent is in charge of two table tennis rackets and their main goal is to act collaboratively and maintain the ball in play for as long as possible, until a certain score is reached.

The project is composed by five files: *Tennis.ipynb* a Jupyter Notebook file containing the main code to initialize dependencies, environment and Agents; *maddpg_agent.py*, a code that contains the characteristics of the multi-agent and how it behaves through this task; *model.py*, containing the deep neural networks (*DNN*) architectures used by the agent; *solved_checkpoint_actor.pth* and *solved_checkpoint_critic.pth* files with saved *DNN*'s weights, that solved the environment.

The problem trying to be solved is modeled as a Markov Decision Process, involving *mappings* from states to actions, called *policy*, in such way that these actions will maximize the total cumulative reward signal of the agent. States are any information that the environment provides the agent, excluding the reward signal. Actions are ways that an agent can interact with the environment, and rewards are signals that the agent receive after interacting with the environment, shaping its learning.

The solution to the problem, on this project, is obtained by utilizing a Policy-Based Method called Multi-Agent Deep Deterministic Policy Gradient (*MADDPG*). Using *DNN*'s as non-linear function approximators, the algorithm can approximate an optimal *policy*. The model take as input a given *state* and outputs the best action to be taken, in that state, to its respective agent.

Implementation

Preparation

The goal is to train an agent able to get an average score of +0.5 over 100 consecutive episodes. The scores are distributed like follows: +0.1 each step an racket can hit the ball over an net and -0.1 if it hit the table or gets thrown out of it's bounds. As the agents interact with the environment, the reward signal guides them towards maintaining the ball in play, characterizing a collaboration scenario.

At first, in the notebook file, the dependencies are installed, libraries are imported, the simulation environment is initialized.

The next step is to explore the State and Action Spaces. The observation space consists of 8 variables corresponding to the position and velocity of the ball and racket. Each agent receives its own, local observation. Two continuous actions are available, corresponding to movement toward (or away from) the net, and jumping.

To learn how the Python API controls the agent and receives the feedbacks from the environment, a code cell is provided with a random action agent.

Learning Algorithm

The *maddpg* algorithm is an approximate Actor-Critic Method, but also resembles the DQN approach of Reinforcement Learning. The agent is composed of two Neural Networks (*NNs*) the Actor and the Critic, both with target and local networks totalizing 4 *NNs*, these are used to encode the *policy* function.

The learning pipeline takes first a state as input in the Actor network, outputting the best possible action in that state, this procedure makes possible for ddpq to tackle continuous action spaces, in contrast to the regular DQN approach. This action is used in the Critic network, alongside with the state, where it outputs an action value function (*q*), this *q* is used as a baseline for updating both Actor and Critic networks, reducing the variance and instability of classic RL algorithms. The optimization is done with a gradient ascent between both Actor's and Critic's target and local networks parameters.

The behaviour of the agent can be explored in the *maddpg_agent.py* file. Important libraries and components are imported and local parameters are initialized: *BUFFER_SIZE*, defines the replay buffer size, shared by the agents, this is an object that contains tuples called experiences composed by state,actions,rewards,next states and dones, these are necessary informations for learning; *BATCH_SIZE*, when the number of experiences in the replay buffer exceeds the batch size, the learning method is called; *TAU*, this hyperparameter controls the model *soft updates*, a method used for

slowly changing the target networks parameters, improving stability; *LR_ACTOR* and *LR_CRITIC*, the optimizer learning rates, these control the gradient ascent step; *WEIGHT_DECAY*, the l2 regularization parameter of the optimizer.

The main implementation begins on fourth step: additional libraries and components are imported, an *agent* is created and initialized with proper parameters: *state_size* and *action_size*. The *maddpg* function is created, taking as parameters the number of episodes (*n_episodes*) and the maximum length of each episode (*max_t*).

In each episode the environment is reseted and the agents receive initial states. While the number of timesteps is less than *max_t*, the following procedures are done:

The agent use it's *act* method with the current state as input, the method takes the input and passes it through the actor network, returning an action for the state. A environment *step* is taken, using the previous obtained action, and it's returns: next state, rewards and dones (if the episode is terminated or not). These are stored in the *env_info* variable, that passes them individually for each of these information's new variables. The agent uses it *step* method, the method first adds the experience tuple for the shared replay buffer and, depending on the size, calls the *learn* method. The rewards are added to the scores variable and the state receives the next state, to give continuation to the environment, if any of the components of the done variable indicates that the episode is over, the loop of *max_t* breaks, and a new episode is initialized.

If the average score of the last 100 episodes is bigger than 0.5, the networks weights are save and the loop of *n_episodes* breaks and the *maddpg* function returns a list with each episode's score. This list is plotted with the episodes number, showing the Agent's learning during the algorithm's execution.

Neural Network Architecture

Actor

Composed of three hidden layers with 100, 75, 75 nodes respectively. The first one with batch normalization and each hidden layer is followed by a ReLU activation function. The output layer is followed by a Tanh function, making possible tackling continuous action spaces. The network take as input the state and outputs the best calculated action for this.

Critic

Composed of two hidden layers with 100+(action space size), 100 nodes respectively. The first one with batch normalization and each hidden layer is followed by a ReLU activation function. The network take as input the state and outputs the action value function for the best action outputted by the Actor network.

Rewards per Episode

Next, an image of the Agent's rewards obtained in each episode until solving the environment.




Figure 1: Rewards per episode and running mean of the solving agent.

Results

Simply applying the *maddpg* algorithm to the environment didn't result in learning. By changing hyperparameters and NN's architectures, the agent started to learn and solved the environment.

A table with used hyperparameters is included below:

Hyperparameter	Value
BUFFER_SIZE	int(1e5)
BATCH_SIZE	128
GAMMA	0.99
TAU	0.002
LR_ACTOR	1.5e-3
LR_CRITIC	1.5e-3
WEIGHT_DECAY	0

Table 2: Hyperparameters used by the best solving agent.

Ideas for Future Work

Multi-Agent Reinforcement Learning is a extremely versatile technique, training agents able to discover complex physical and communicative coordination strategies is a marvelous tool. An idea for future work is to create and evaluate a multi-agent environment that resembles a subway, where each agent (train) receives the same observations and the model learns to control all the trains while raising the passenger/time efficiency.